

CYO Machine Learning Project: Titanic - Machine Learning from Disaster

Eddwin Cheteni

2021-03-19

Overview

The data has been split into two groups:

- training set (train.csv)
- test set (test.csv)

The training set will be used to build machine learning models. For the training set, the outcomes are provided (also known as the “ground truth”) for each passenger. The model will be based on “features” like passengers’ gender and class. The test set will be used to see how well our model performs on unseen data. For the test set, the ground truth is not provided for each passenger. The objective of this project is to predict these outcomes. For each passenger in the test set, we use the trained model to predict whether or not they survived the sinking of the Titanic.

The datasets also include gender_submission.csv, a set of predictions that assume all and only female passengers survive, as an example of what a submission file should look like.

Loading data

```
#Loading train set
train_set <- read_csv("~/Data Science Projects/train.csv")
#converting the column names into lowercase
names(train_set) <- tolower(names(train_set))

#Loading test set
test_set <- read_csv("~/Data Science Projects/test.csv")
#converting the column names into lowercase
names(test_set) <- tolower(names(test_set))
```

Data Dictionary

```
## [1] "Dimensions of the train set are:"
## [1] 891  12
## [1] "Dimensions of the test set are:"
## [1] 418  11
```

This confirms that the train set has 891 rows and 12 columns while the test set has 418 rows and 11 columns.

```
## # A tibble: 6 x 12
##   passengerid survived pclass name
sex      age sibsp parch ticket      fare cabin embarked
##      <dbl>      <dbl> <dbl> <chr>
<chr> <dbl> <dbl> <dbl> <chr>      <dbl> <chr> <chr>
## 1          1          0      3 Braund, Mr. Owen Harris
male      22          1      0 A/5 21171      7.25 <NA> S
## 2          2          1      1 Cumings, Mrs. John Bradley (Florence Briggs Thayer)
female    38          1      0 PC 17599      71.3 C85 C
## 3          3          1      3 Heikkinen, Miss. Laina
female    26          0      0 STON/O2. 3101282 7.92 <NA> S
## 4          4          1      1 Futrelle, Mrs. Jacques Heath (Lily May Peel)
female    35          1      0 113803      53.1 C123 S
## 5          5          0      3 Allen, Mr. William Henry
male      35          0      0 373450      8.05 <NA> S
## 6          6          0      3 Moran, Mr. James
male      NA          0      0 330877      8.46 <NA> Q

## # A tibble: 6 x 11
##   passengerid pclass name              sex      age
sibsp parch ticket  fare cabin embarked
##      <dbl>      <dbl> <chr>
<dbl> <dbl> <chr> <dbl> <chr> <chr>
## 1          892          3 Kelly, Mr. James      male     34.5
0      0 330911 7.83 <NA> Q
## 2          893          3 Wilkes, Mrs. James (Ellen Needs)  female    47
1      0 363272 7      <NA> S
## 3          894          2 Myles, Mr. Thomas Francis      male     62
0      0 240276 9.69 <NA> Q
## 4          895          3 Wirz, Mr. Albert      male     27
0      0 315154 8.66 <NA> S
## 5          896          3 Hirvonen, Mrs. Alexander (Helga E Lindqvist) female    22
1      1 3101298 12.3 <NA> S
## 6          897          3 Svensson, Mr. Johan Cervin      male     14
0      0 7538 9.22 <NA> S
```

Data Exploratory Analysis

Visualize train_set

We need to explore on our data so as to get read of unnecessary features and some anomalies in the data such as missing values and outliers.

```
## spec_tbl_df [891 x 12] (S3: spec_tbl_df/tbl_df/tbl/data.frame)
## $ passengerid: num [1:891] 1 2 3 4 5 6 7 8 9 10 ...
## $ survived   : num [1:891] 0 1 1 1 0 0 0 0 1 1 ...
## $ pclass     : num [1:891] 3 1 3 1 3 3 1 3 3 2 ...
## $ name       : chr [1:891] "Braund, Mr. Owen Harris" "Cumings, Mrs. John Bradley
(Florence Briggs Thayer)" "Heikkinen, Miss. Laina" "Futrelle, Mrs. Jacques Heath
(Lily May Peel)" ...
## $ sex        : chr [1:891] "male" "female" "female" "female" ...
## $ age        : num [1:891] 22 38 26 35 35 NA 54 2 27 14 ...
## $ sibsp      : num [1:891] 1 1 0 1 0 0 0 3 0 1 ...
## $ parch      : num [1:891] 0 0 0 0 0 0 0 1 2 0 ...
## $ ticket     : chr [1:891] "A/5 21171" "PC 17599" "STON/O2. 3101282" "113803" ...
```

```
## $ fare      : num [1:891] 7.25 71.28 7.92 53.1 8.05 ...
## $ cabin     : chr [1:891] NA "C85" NA "C123" ...
## $ embarked  : chr [1:891] "S" "C" "S" "S" ...
## - attr(*, "spec")=
## .. cols(
## .. PassengerId = col_double(),
## .. Survived = col_double(),
## .. Pclass = col_double(),
## .. Name = col_character(),
## .. Sex = col_character(),
## .. Age = col_double(),
## .. SibSp = col_double(),
## .. Parch = col_double(),
## .. Ticket = col_character(),
## .. Fare = col_double(),
## .. Cabin = col_character(),
## .. Embarked = col_character()
## .. )
```

Lets have a look at the summary of our train set:


```
##  passengerid    survived      pclass      name      sex
age      sibsp      parch      ticket      fare
## Min.   : 1      Min.   :0.000      Min.   :1.00      Length:891      Length:891
Min.   : 0.42      Min.   :0.000      Min.   :0.000      Length:891      Min.   : 0.00
## 1st Qu.:224      1st Qu.:0.000      1st Qu.:2.00      Class :character      Class :character
1st Qu.:20.12      1st Qu.:0.000      1st Qu.:0.000      Class :character      1st Qu.: 7.91
## Median :446      Median :0.000      Median :3.00      Mode  :character      Mode  :character
Median :28.00      Median :0.000      Median :0.000      Mode  :character      Median : 14.45
## Mean   :446      Mean   :0.384      Mean   :2.31
Mean   :29.70      Mean   :0.523      Mean   :0.382
## 3rd Qu.:668      3rd Qu.:1.000      3rd Qu.:3.00
3rd Qu.:38.00      3rd Qu.:1.000      3rd Qu.:0.000
## Max.   :891      Max.   :1.000      Max.   :3.00
Max.   :80.00      Max.   :8.000      Max.   :6.000
##
NA's    :177
##  cabin      embarked
## Length:891      Length:891
## Class :character      Class :character
## Mode  :character      Mode  :character
##
##
##
##
```

It indicates that age has 177 missing values.

```
#checking variables with missing values
list_na <- colnames(train_set)[ apply(train_set, 2, anyNA) ]
list_na

## [1] "age"      "cabin"    "embarked"
```

```
#checking missing values by use of a barplot on embarked feature
train_set%>%
  select(embarked)%>% #select embarked column
  ggplot()+
  geom_bar(aes(embarked), stat = "count", na.rm=TRUE, position = "stack", fill
=I("tomato3"))+
  ggtitle("Bar plot for Embarked Feature") #title of the chart
```



Based on the data source provided, there is 20% missing values for age, 77% for cabin and almost 0% for embarked. From this observation, we can drop the cabin column, replace with median on the age column and mode in the embarked column.

```
#since Southampton has the most passengers
embarked_mode<-"S"
#computing the median age
age_median<-median(train_set$age,na.rm = TRUE)
#replacing the missing values, age by median and embarked by the most frequent one
i.e. Southampton
train_set_replaced <- train_set %>%
  mutate(age = ifelse(is.na(age), age_median, age),embarked =
ifelse(is.na(embarked),embarked_mode, embarked))
head(train_set_replaced)

## # A tibble: 6 x 12
##   passengerid survived pclass name      fare cabin embarked
##       <dbl>      <dbl>   <dbl> <chr>    <dbl> <chr>   <chr>
## 1             1         0     3 Braund, Mr. Owen Harris
```

```

male      22      1      0 A/5 21171      7.25 <NA> S
## 2      2      1      1 Cumings, Mrs. John Bradley (Florence Briggs Thayer)
female    38      1      0 PC 17599      71.3 C85 C
## 3      3      1      3 Heikkinen, Miss. Laina
female    26      0      0 STON/O2. 3101282 7.92 <NA> S
## 4      4      1      1 Futrelle, Mrs. Jacques Heath (Lily May Peel)
female    35      1      0 113803      53.1 C123 S
## 5      5      0      3 Allen, Mr. William Henry
male      35      0      0 373450      8.05 <NA> S
## 6      6      0      3 Moran, Mr. James
male      28      0      0 330877      8.46 <NA> Q

```

Dropping unnecessary columns

The variable column Cabin is dropped since 77% of the observations are missing therefore, including this in our model might compromise our model performance.

```

colnames(train_set_replaced)[apply(train_set_replaced, 2, anyNA)]

## [1] "cabin"

# Data pre-processing on train set
df<- train_set_replaced %>%
  mutate(survived = as.factor(survived),#convert to a factor
         pclass = as.factor(pclass),#convert to a factor
         sex = as.factor(sex),#convert to a factor
         embarked = as.factor(embarked),#convert to a factor
         fare = log1p(fare),#log-transform fare variable
         nbr_family = as.integer(sibsp + parch),#combine siblings and parents
         together
         age = scale(age)) %>%#convert to an integer
  select(survived, pclass, sex, age, nbr_family, fare, embarked)#selecting the
required columns
# Add cleaner factor levels
levels(df$survived) <- c('Perished','Survived')
levels(df$embarked) <- c('Cherbourg','Queenstown','Southampton')
levels(df$sex) <- c('Female','Male')
df%>%head()

## # A tibble: 6 x 7
##   survived pclass sex   age[,1] nbr_family fare embarked
##   <fct>    <fct> <fct>   <dbl>    <int> <dbl> <fct>
## 1 Perished 3      Male  -0.565      1  2.11 Southampton
## 2 Survived 1      Female 0.663      1  4.28 Cherbourg
## 3 Survived 3      Female -0.258      0  2.19 Southampton
## 4 Survived 1      Female 0.433      1  3.99 Southampton
## 5 Perished 3      Male   0.433      0  2.20 Southampton
## 6 Perished 3      Male  -0.105      0  2.25 Queenstown

df%>%str()

## tibble [891 x 7] (S3: tbl_df/tbl/data.frame)
## $ survived : Factor w/ 2 levels "Perished","Survived": 1 2 2 2 1 1 1 1 2 2 ...
## $ pclass : Factor w/ 3 levels "1","2","3": 3 1 3 1 3 3 1 3 3 2 ...
## $ sex : Factor w/ 2 levels "Female","Male": 2 1 1 1 2 2 2 2 1 1 ...

```

```
## $ age      : num [1:891, 1] -0.565 0.663 -0.258 0.433 0.433 ...
##   ..- attr(*, "scaled:center")= num 29.4
##   ..- attr(*, "scaled:scale")= num 13
## $ nbr_family: int [1:891] 1 1 0 1 0 0 0 4 2 1 ...
## $ fare      : num [1:891] 2.11 4.28 2.19 3.99 2.2 ...
## $ embarked  : Factor w/ 3 levels "Cherbourg","Queenstown",...: 3 1 3 3 3 2 3 3 3 1
...
```

Exploring the test set:

```
#checking missing values on test set
colnames(test_set)[ apply(test_set, 2, anyNA) ]

## [1] "age" "fare" "cabin"

#computing the median age
age_median<-median(test_set$age,na.rm = TRUE)
#computing the median fare
fare_median<-median(test_set$fare,na.rm = TRUE)
#replacing the missing values in age and fare column so that validation can be performed
test_set_replaced <- test_set %>%
  mutate(age = ifelse(is.na(age), age_median, age),fare =
ifelse(is.na(fare),fare_median, fare))
head(test_set_replaced)

## # A tibble: 6 x 11
##   passengerid pclass name          sex    age
##   sibsp parch ticket   fare cabin embarked
##   <dbl> <dbl> <dbl> <dbl> <chr>   <chr> <dbl>
## 1      892     3 Kelly, Mr. James  male   34.5
0      0 330911  7.83 <NA> Q
## 2      893     3 Wilkes, Mrs. James (Ellen Needs) female  47
1      0 363272  7     <NA> S
## 3      894     2 Myles, Mr. Thomas Francis  male   62
0      0 240276  9.69 <NA> Q
## 4      895     3 Wirz, Mr. Albert      male   27
0      0 315154  8.66 <NA> S
## 5      896     3 Hirvonen, Mrs. Alexander (Helga E Lindqvist) female  22
1      1 3101298 12.3  <NA> S
## 6      897     3 Svensson, Mr. Johan Cervin  male   14
0      0 7538    9.22 <NA> S

#checking missing values on test set
colnames(test_set_replaced)[ apply(test_set_replaced, 2, anyNA) ]

## [1] "cabin"
```

Let us create the validation set from the test set:

```
# Data pre-processing on test set to create the validation set
validation<- test_set_replaced %>%
  mutate(pclass = as.factor(pclass),#convert to a factor
sex = as.factor(sex),#convert to a factor
embarked = as.factor(embarked),#convert to a factor
```

```

    fare = log1p(fare),#Log-transform variable fare
    nbr_family = as.integer(sibsp + parch),#combine siblings and parents
together
    age = scale(age)) %>% #Normalize age
    select(pclass, sex, age, nbr_family, fare, embarked)#selecting important features
# Add cleaner factor levels
levels(validation$embarked) <- c('Cherbourg','Queenstown','Southampton')
levels(validation$sex) <- c('Female','Male')
validation%>%head()

## # A tibble: 6 x 6
##   pclass sex    age[,1] nbr_family  fare embarked
##   <fct> <fct>    <dbl>      <int> <dbl> <fct>
## 1 3      Male    0.386         0  2.18 Queenstown
## 2 3      Female  1.37         1  2.08 Southampton
## 3 2      Male    2.55         0  2.37 Queenstown
## 4 3      Male   -0.205        0  2.27 Southampton
## 5 3      Female -0.598         2  2.59 Southampton
## 6 3      Male   -1.23         0  2.32 Southampton

#####
# Pre-Processing
#####

#Create dummy variables on the categorical features in df set
df <- dummy_cols(df,
                  select_columns=c('pclass','sex', 'embarked'),
                  remove_most_frequent_dummy = TRUE) %>%
  select(-pclass,-sex,-embarked)
#Create dummy variables on the categorical features in validation set
validation <- dummy_cols(validation,
                          select_columns=c('pclass','sex', 'embarked'),
                          remove_most_frequent_dummy = TRUE) %>%
  select(-pclass,-sex,-embarked)

```

The tables below shows the datasets (df and validation set) ready for model training process after dummies were created in the above procedure.

```

#Have a Look at the df set ready for splitting
df%>%head()%>%knitr::kable()

```

survived	age	nbr_fami ly	fare	pclas s_1	pclas s_2	sex_Fema le	embarked_C herbourg	embarked _Queenst own
Perished	-0.56542	1	2.1102	0	0	0	0	0
Survived	0.66349	1	4.2806	1	0	1	1	0
Survived	-0.25819	0	2.1889	0	0	1	0	0
Survived	0.43307	1	3.9908	1	0	1	0	0
Perished	0.43307	0	2.2028	0	0	0	0	0
Perished	-0.10458	0	2.2469	0	0	0	0	1

```
#Have a look at the validation set ready for final model validation
validation%>%head()%>%knitr::kable()
```

age	nbr_fam ily	fare	pclass _1	pclass _2	sex_Fe male	embarked_ Cherbourg	embarked_ Queenstow n
0.38577	0	2.1781	0	0	0	0	1
1.36973	1	2.0794	0	0	1	0	0
2.55048	0	2.3691	0	1	0	0	1
-0.20461	0	2.2683	0	0	0	0	0
-0.59819	2	2.5868	0	0	1	0	0
-1.22793	0	2.3248	0	0	0	0	0

Data partition

We create this partition on the df set so that the validation set will be used for validation of the model.

```
## [1] "Dimensions of the train set are:"
## [1] 714 9
## [1] "Dimensions of the test set are:"
## [1] 177 9
```

Model building

Model 1: Logistic Regression Model (glm)

```
#####
# Model 1: Logistic Regression Model ('glm') all selected variables
#####

# Fit model with all variables
(fit_glm <- glm(survived ~ .,
               data=train,
               family=binomial))

##
## Call: glm(formula = survived ~ ., family = binomial, data = train)
##
## Coefficients:
## (Intercept)          age      nbr_family          fare
pclass_1      pclass_2    sex_Female  embarked_Ch
##          -3.303      -0.519      -0.288          0.405
1.690          1.107          2.749          0.338
## embarked_Queenstown
##          0.184
##
```



```
## Degrees of Freedom: 713 Total (i.e. Null); 705 Residual
## Null Deviance: 951
## Residual Deviance: 617 AIC: 635

# Generate prediction
p_hat_glm <- predict(fit_glm, test, type='response')
y_hat_glm <- ifelse(p_hat_glm > 0.5, "Survived", "Perished") %>% factor()
# Compute the accuracy
acc <- confusionMatrix(y_hat_glm, test$survived, positive='Survived')
#Generate table of accuracy
accuracy_results <- tibble(method='Model 1: Logistic Regression Model',
                           accuracy = acc$overall['Accuracy'])
accuracy_results%>%knitr::kable()
```

method	accuracy
Model 1: Logistic Regression Model	0.78531

Model 2: Random Forest Model (rf)

```
#####
# Model 2: Random Forest Model ('rf')
#####

set.seed(1234)

# Set control parameters
control <- trainControl(method='repeatedcv',
                        number=10,
                        repeats=5,
                        search = 'random')

# Determine baseline mtry
mtry <- sqrt(ncol(train))
tuneGrid = expand.grid(.mtry=mtry)

#Train RF model
#Random generate 15 mtry values with tuneLength = 15
train_rf <- train(survived ~ .,
                 data=train,
                 method='rf',
                 tuneLength=15,
                 trControl=control,
                 importance=TRUE,
                 localImp=TRUE)

# Explain final RF model
(fit_rf <- train_rf$finalModel)

##
## Call:
## randomForest(x = x, y = y, mtry = param$mtry, importance = TRUE, localImp =
TRUE)
##
## Type of random forest: classification
## Number of trees: 500
```

```

## No. of variables tried at each split: 4
##
##          OOB estimate of  error rate: 16.67%
## Confusion matrix:
##          Perished Survived class.error
## Perished      388      52      0.11818
## Survived       67     207      0.24453

#Generate predictions
y_hat_rf <- predict(fit_rf, test)
#Compute the accuracy
acc <- confusionMatrix(y_hat_rf, test$survived, positive='Survived')
acc

## Confusion Matrix and Statistics
##
##          Reference
## Prediction Perished Survived
##   Perished      94      20
##   Survived      15      48
##
##          Accuracy : 0.802
##          95% CI : (0.736, 0.858)
##   No Information Rate : 0.616
##   P-Value [Acc > NIR] : 7.43e-08
##
##          Kappa : 0.576
##
##  Mcnemar's Test P-Value : 0.499
##
##          Sensitivity : 0.706
##          Specificity : 0.862
##          Pos Pred Value : 0.762
##          Neg Pred Value : 0.825
##          Prevalence : 0.384
##          Detection Rate : 0.271
##          Detection Prevalence : 0.356
##          Balanced Accuracy : 0.784
##
##          'Positive' Class : Survived
##

#Generate table for results
accuracy_results <- accuracy_results %>%
  bind_rows(tibble(method='Model 2: Random Forest Model',
                    accuracy = acc$overall['Accuracy'])))
accuracy_results %>% knitr::kable()

```

method	accuracy
Model 1: Logistic Regression Model	0.78531
Model 2: Random Forest Model	0.80226

Model 3: Gradient Boosting Model (gbm)

```
#####  
# Model 3: Gradient Boosting Model ('gbm')  
#####  
  
set.seed(1234)  
  
# Set control parameters  
control <- trainControl(method='repeatedcv',  
                        number=4,  
                        repeats=4)  
  
#Train gbm model  
fit_gbm <- train(survived ~ .,  
                data=train,  
                method='gbm',  
                trControl=control,  
                verbose = FALSE)  
  
fit_gbm  
  
## Stochastic Gradient Boosting  
##  
## 714 samples  
## 8 predictor  
## 2 classes: 'Perished', 'Survived'  
##  
## No pre-processing  
## Resampling: Cross-Validated (4 fold, repeated 4 times)  
## Summary of sample sizes: 536, 536, 535, 535, 535, 536, ...  
## Resampling results across tuning parameters:  
##  
## interaction.depth n.trees Accuracy Kappa  
## 1 50 0.80917 0.58808  
## 1 100 0.81234 0.59836  
## 1 150 0.80955 0.59232  
## 2 50 0.81619 0.60545  
## 2 100 0.81514 0.60403  
## 2 150 0.81619 0.60663  
## 3 50 0.81934 0.61134  
## 3 100 0.82423 0.62067  
## 3 150 0.82353 0.61979  
##  
## Tuning parameter 'shrinkage' was held constant at a value of 0.1  
## Tuning parameter 'n.minobsinnode' was held constant at a value of 10  
## Accuracy was used to select the optimal model using the largest value.  
## The final values used for the model were n.trees = 100, interaction.depth = 3,  
shrinkage = 0.1 and n.minobsinnode = 10.  
  
#Generate predictions  
y_hat_gbm <- predict(fit_gbm, test, type = "raw")  
  
#Compute the accuracy
```

```

acc <- confusionMatrix(y_hat_gbm, test$survived, positive='Survived')
acc

## Confusion Matrix and Statistics
##
##              Reference
## Prediction Perished Survived
##   Perished      96      24
##   Survived      13      44
##
##              Accuracy : 0.791
##              95% CI : (0.724, 0.848)
##   No Information Rate : 0.616
##   P-Value [Acc > NIR] : 4.55e-07
##
##              Kappa : 0.544
##
##   Mcnemar's Test P-Value : 0.1
##
##              Sensitivity : 0.647
##              Specificity : 0.881
##   Pos Pred Value : 0.772
##   Neg Pred Value : 0.800
##   Prevalence : 0.384
##   Detection Rate : 0.249
##   Detection Prevalence : 0.322
##   Balanced Accuracy : 0.764
##
##   'Positive' Class : Survived
##

#Generate table for results
accuracy_results <- accuracy_results %>%
  bind_rows(tibble(method='Model 3: Gradient Boosting Model',
                    accuracy = acc$overall['Accuracy']))
accuracy_results %>% knitr::kable()

```

method	accuracy
Model 1: Logistic Regression Model	0.78531
Model 2: Random Forest Model	0.80226
Model 3: Gradient Boosting Model	0.79096

Model 4: Support Vector Machine Model (svm)

```

#####
# Model 4: Support Vector Machine Model ('svm')
#####
set.seed(1234)

#setting control parameters
cv_opts = trainControl(method="cv", number=10)

#Train the model

```

```

results_svm = train(survived~.,
                    data=train,
                    method="svmLinear",
                    preProcess="range",
                    trControl=cv_opts,
                    tuneLength=5)

results_svm

## Support Vector Machines with Linear Kernel
##
## 714 samples
## 8 predictor
## 2 classes: 'Perished', 'Survived'
##
## Pre-processing: re-scaling to [0, 1] (8)
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 643, 643, 642, 643, 643, 643, ...
## Resampling results:
##
## Accuracy Kappa
## 0.7928 0.5562
##
## Tuning parameter 'C' was held constant at a value of 1

#Generate predictions
y_hat_svm = predict(results_svm, test)
#Compute the accuracy
acc <- confusionMatrix(y_hat_svm, test$survived, positive='Survived')
acc

## Confusion Matrix and Statistics
##
##              Reference
## Prediction Perished Survived
## Perished      92      25
## Survived      17      43
##
##              Accuracy : 0.763
##              95% CI : (0.693, 0.823)
## No Information Rate : 0.616
## P-Value [Acc > NIR] : 2.39e-05
##
##              Kappa : 0.487
##
## Mcnemar's Test P-Value : 0.28
##
##              Sensitivity : 0.632
##              Specificity : 0.844
##              Pos Pred Value : 0.717
##              Neg Pred Value : 0.786
##              Prevalence : 0.384
##              Detection Rate : 0.243
##              Detection Prevalence : 0.339
##              Balanced Accuracy : 0.738
##

```

```
##      'Positive' Class : Survived
##

#Generate table for results
accuracy_results <- accuracy_results %>%
  bind_rows(tibble(method='Model 4: Support Vector Machine Model',
                    accuracy = acc$overall['Accuracy']))
accuracy_results %>% knitr::kable()
```

method	accuracy
Model 1: Logistic Regression Model	0.78531
Model 2: Random Forest Model	0.80226
Model 3: Gradient Boosting Model	0.79096
Model 4: Support Vector Machine Model	0.76271

Model 5: Neural Network Model (nnet)

```
#####
# Model 5: Neural Network Model ('nnet')
#####
results_nnet = train(survived~.,
                     data=train,
                     method="avNNet",
                     trControl=cv_opts,
                     preProcess="range",
                     tuneLength=5,
                     trace=F,
                     maxit=1000)

results_nnet

## Model Averaged Neural Network
##
## 714 samples
## 8 predictor
## 2 classes: 'Perished', 'Survived'
##
## Pre-processing: re-scaling to [0, 1] (8)
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 643, 643, 643, 642, 642, ...
## Resampling results across tuning parameters:
##
##  size  decay  Accuracy  Kappa
##  1      0e+00  0.80379   0.58087
##  1      1e-04  0.79832   0.56023
##  1      1e-03  0.80115   0.56106
##  1      1e-02  0.80395   0.56580
##  1      1e-01  0.80256   0.56428
##  3      0e+00  0.82349   0.61770
##  3      1e-04  0.82079   0.60906
##  3      1e-03  0.81230   0.59284
##  3      1e-02  0.81939   0.60990
##  3      1e-01  0.80677   0.57740
##  5      0e+00  0.82218   0.61347
```

```

## 5      1e-04  0.82218  0.61836
## 5      1e-03  0.82494  0.62700
## 5      1e-02  0.81379  0.60170
## 5      1e-01  0.80675  0.57662
## 7      0e+00  0.81800  0.60660
## 7      1e-04  0.80266  0.58016
## 7      1e-03  0.80959  0.59288
## 7      1e-02  0.80959  0.59284
## 7      1e-01  0.80816  0.58010
## 9      0e+00  0.82076  0.61236
## 9      1e-04  0.81097  0.59699
## 9      1e-03  0.80117  0.57996
## 9      1e-02  0.80115  0.57520
## 9      1e-01  0.80816  0.58010
##
## Tuning parameter 'bag' was held constant at a value of FALSE
## Accuracy was used to select the optimal model using the largest value.
## The final values used for the model were size = 5, decay = 0.001 and bag = FALSE.

#Generate predictions
y_hat_nnet <- predict(results_nnet, test)
#Compute the accuracy
acc <- confusionMatrix(y_hat_nnet, test$survived, positive="Survived")
acc

## Confusion Matrix and Statistics
##
##              Reference
## Prediction Perished Survived
##   Perished      97      21
##   Survived      12      47
##
##              Accuracy : 0.814
##              95% CI : (0.748, 0.868)
##   No Information Rate : 0.616
##   P-Value [Acc > NIR] : 1.06e-08
##
##              Kappa : 0.596
##
## Mcnemar's Test P-Value : 0.164
##
##              Sensitivity : 0.691
##              Specificity : 0.890
##   Pos Pred Value : 0.797
##   Neg Pred Value : 0.822
##   Prevalence : 0.384
##   Detection Rate : 0.266
##   Detection Prevalence : 0.333
##   Balanced Accuracy : 0.791
##
##   'Positive' Class : Survived
##

#Generate table for results
accuracy_results <- accuracy_results %>%

```

```

    bind_rows(tibble(method='Model 5: Neural Network Model',
                      accuracy = acc$overall['Accuracy']))
accuracy_results %>% knitr::kable()

```

method	accuracy
Model 1: Logistic Regression Model	0.78531
Model 2: Random Forest Model	0.80226
Model 3: Gradient Boosting Model	0.79096
Model 4: Support Vector Machine Model	0.76271
Model 5: Neural Network Model	0.81356

Neural Network model gives the best accuracy of 0.81356 on the test set, so we will consider it in predicting survivors in the validation set to be submitted on Kaggle site for evaluation.

Final model - Neural network (nnet)

```

#####
# Final Model
# Based on modeled results, apply Neural Network Model
# to 'test set' and create submission csv
#####
#setting control parameters
cv_opts = trainControl(method="cv", number=10)

#Train nnet model
fit_final <- train(survived~.,
                  data=train,
                  method="avNNet",
                  trControl=cv_opts,
                  preProcess="range",
                  tuneLength=5,
                  trace=F,
                  maxit=1000)

fit_final

## Model Averaged Neural Network
##
## 714 samples
## 8 predictor
## 2 classes: 'Perished', 'Survived'
##
## Pre-processing: re-scaling to [0, 1] (8)
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 642, 643, 643, 643, 643, 642, ...
## Resampling results across tuning parameters:
##
##  size  decay  Accuracy  Kappa
##  1      0e+00  0.80115   0.56649
##  1      1e-04  0.80822   0.57847
##  1      1e-03  0.80962   0.57792
##  1      1e-02  0.80962   0.57695
##  1      1e-01  0.80964   0.58004

```



```
## 3 0e+00 0.82367 0.61373
## 3 1e-04 0.81800 0.60558
## 3 1e-03 0.81941 0.60496
## 3 1e-02 0.82226 0.61135
## 3 1e-01 0.81377 0.58771
## 5 0e+00 0.81794 0.60423
## 5 1e-04 0.82222 0.61495
## 5 1e-03 0.81800 0.60842
## 5 1e-02 0.81516 0.59970
## 5 1e-01 0.81659 0.59472
## 7 0e+00 0.81939 0.61117
## 7 1e-04 0.81377 0.60043
## 7 1e-03 0.82502 0.62356
## 7 1e-02 0.81657 0.60372
## 7 1e-01 0.81659 0.59472
## 9 0e+00 0.81234 0.59305
## 9 1e-04 0.80953 0.59309
## 9 1e-03 0.81516 0.60234
## 9 1e-02 0.81520 0.60177
## 9 1e-01 0.81659 0.59472
##
## Tuning parameter 'bag' was held constant at a value of FALSE
## Accuracy was used to select the optimal model using the largest value.
## The final values used for the model were size = 7, decay = 0.001 and bag = FALSE.

#Generate predictions
final_pred <- predict(fit_final, validation)
# Save the solution to a dataframe with two columns: PassengerId and Survived
(prediction)
gender_submission <- read_csv("~/Data Science Projects/gender_submission.csv")
results <- data.frame(PassengerID = gender_submission$PassengerId)
solution <- results %>% mutate(Survived = ifelse(final_pred=="Survived",1,0))
# Write the solution to file
write_csv(solution, path = '~/Data Science Projects/final_solution.csv', col_names =
TRUE)
```

Final Model accuracy

```
actual<-gender_submission%>%mutate(outcome = ifelse(gender_submission$Survived ==
'1', "Survived", "Perished"))
predicted <- solution%>%mutate(outcome = ifelse(solution$Survived == '1', "Survived",
"Perished"))
#Compute the accuracy
acc <- confusionMatrix(as.factor(predicted$outcome), as.factor(actual$outcome),
positive="Survived")
acc

## Confusion Matrix and Statistics
##
##           Reference
## Prediction Perished Survived
##   Perished      231      29
##   Survived       35     123
##
##           Accuracy : 0.847
```

```

##           95% CI : (0.809, 0.88)
##   No Information Rate : 0.636
##   P-Value [Acc > NIR] : <2e-16
##
##           Kappa : 0.672
##
##   McNemar's Test P-Value : 0.532
##
##           Sensitivity : 0.809
##           Specificity : 0.868
##           Pos Pred Value : 0.778
##           Neg Pred Value : 0.888
##           Prevalence : 0.364
##           Detection Rate : 0.294
##           Detection Prevalence : 0.378
##           Balanced Accuracy : 0.839
##
##           'Positive' Class : Survived
##
#Generate table for results
final_accuracy <- acc$overall['Accuracy']
final_accuracy

## Accuracy
## 0.84689

```

Conclusion

Looking at the outcome obtained above, an accuracy of 0.84689 was obtained when computed on the predicted survivors in the provided gender_submission dataset against the predicted survivors using the nnet model. The final model results (final_solution file) were submitted on <https://www.kaggle.com/submissions/20054628/20054628.raw> for evaluation and scored an overall accuracy of 0.78468 on an unseen data, which sounds to be a good model.

References

1. *Titanic - Machine Learning from Disaster* <https://www.kaggle.com/c/titanic>
2. Clark, M. (2013), *"An Introduction to Machine Learning with Application in R"*, Center for Social Research, University of Notre Dame.